

Interactive Computer Graphics: Lecture 9

Rasterization, Visibility & Anti-aliasing

Some slides adopted from
F. Durand and B. Cutler, MIT

The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display

- Rasterizes objects into pixels
- Interpolate values inside objects (color, depth, etc.)



Graphics Lecture 9: Slide 2

The Graphics Pipeline

Modelling
Transformations

Illumination
(Shading)

Viewing Transformation
(Perspective / Orthographic)


Clipping

Projection
(to Screen Space)

Scan Conversion
(Rasterization)

Visibility / Display


- Handles occlusions
- Determines which objects are closest and therefore visible



Graphics Lecture 9: Slide 3

Rasterization


- Determine which pixels are drawn into the framebuffer
- Interpolate parameters (colors, texture coordinates, etc.)



Graphics Lecture 9: Slide 4


Rasterization

- What does interpolation mean?
- Examples: Colors, normals, shading, texture coordinates



Graphics Lecture 9: Slide 5


Coordinate intuition



Graphics Lecture 9: Slide 6

Coordinate intuition


Trilinear coordinates



Graphics Lecture 9: Slide 7

Coordinate intuition


barycentric coordinates



Graphics Lecture 9: Slide 8

Coordinate intuition


barycentric coordinates



Graphics Lecture 9: Slide 9

A triangle in terms of vectors


- We can use vertices **a**, **b** and **c** to specify the three points of a triangle
- We can also compute the edge vectors



Graphics Lecture 9: Slide 10

Points and planes

- The three non-collinear points determine a plane




- Example: The vertices a , b and c determine a plane
- The vectors $b - a$ and $c - a$ form a basis for this plane

Basis vectors

- This (non-orthogonal) basis can be used to specify the location of any point **p** in the plane

$$\mathbf{p} = \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a})$$



Barycentric coordinates

- We can reorder the terms of the equation:

$$\begin{aligned}\mathbf{p} &= \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a}) \\ &= (1 - \beta - \gamma)\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c} \\ &= \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}\end{aligned}$$

- In other words:

$$\mathbf{p}(\alpha, \beta, \gamma) = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

- α , β , γ and called barycentric coordinates

Barycentric coordinates

- **Homogenous barycentric coordinates:**
 - normalised so that $\alpha + \beta + \gamma = \text{area of triangle}$
- **Areal coordinates or absolute barycentric coordinates** : barycentric coordinates *normalized by the area of the original triangle* $\alpha + \beta + \gamma = 1$

Barycentric coordinates

- Barycentric coordinates describe a point **p** as an affine combination of the triangle vertices

$$\mathbf{p}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c} \quad \alpha + \beta + \gamma = 1$$

- For any point **p** inside the triangle (**a**, **b**, **c**):

$$0 < \alpha < 1$$


$$0 < \beta < 1$$

$$0 < \gamma < 1$$

- Point on an edge: one coefficient is 0
- Vertex: two coefficients are 0, remaining one is 1

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 16

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 17

Barycentric coordinates and signed distances


- Let $\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$. Each coordinate (e.g. β) is the signed distance from \mathbf{p} to the line through a triangle edge (e.g. \mathbf{ac})



Graphics Lecture 9: Slide 18

Barycentric coordinates and signed distances

- The signed distance can be computed by evaluating implicit line equations, e.g., $f_{ac}(x,y)$ of edge **ac**



Graphics Lecture 9: Slide 19

Recall: Implicit equation for lines

- Implicit equation in 2D:

$$f(x, y) = 0$$

- Points with $f(x, y) = 0$ are on the line
- Points with $f(x, y) \neq 0$ are not on the line

- General implicit form

$$Ax + By + C = 0$$

- Implicit line through two points (x_a, y_a) and (x_b, y_b)


$$(y_a - y_b)x + (x_b - x_a)y + x_a y_b - x_b y_a = 0$$

Implicit equation for lines: Example

A =

B =

C =




Graphics Lecture 9: Slide 21

Implicit equation for lines: Example

Solution 1: $-2x + 4y = 0$

Solution 2: $2x - 4y = 0$

$$kf(x, y) = 0 \text{ for any } k$$



Graphics Lecture 9: Slide 22


Edge equations

- Given a triangle with vertices (x_a, y_a) , (x_b, y_b) , and (x_c, y_c) .
- The line equations of the edges of the triangle are:

$$f_{ab}(x, y) = (y_a - y_b)x + (x_b - x_a)y + x_a y_b - x_b y_a$$

$$f_{bc}(x, y) = (y_b - y_c)x + (x_c - x_b)y + x_b y_c - x_c y_b$$

$$f_{ca}(x, y) = (y_c - y_a)x + (x_a - x_c)y + x_c y_a - x_a y_c$$



Graphics Lecture 9: Slide 23

Barycentric Coordinates

- Remember that: $f(x,y) = 0 \Leftrightarrow kf(x,y) = 0$
- A barycentric coordinate (e.g. β) is a signed distance from a line (e.g. the line that goes through **ac**)
- For a given point **p**, we would like to compute its barycentric coordinate β using an implicit edge equation.
- We need to choose k such that

$$kf_{ac}(x,y) = \beta$$

Barycentric Coordinates

- We would like to choose k such that: $kf_{ac}(x,y) = \beta$
- We know that $\beta = 1$ at point \mathbf{b} :

$$kf_{ac}(x,y) = 1 \Leftrightarrow k = \frac{1}{f_{ac}(x_b, y_b)}$$

- The barycentric coordinate β for point \mathbf{p} is:

$$\beta = \frac{f_{ac}(x,y)}{f_{ac}(x_b, y_b)}$$

Barycentric Coordinates

- In general, the barycentric area coordinates for point **p** are:


$$\alpha = \frac{f_{bc}(x, y)}{f_{bc}(x_a, y_a)} \quad \beta = \frac{f_{ac}(x, y)}{f_{ac}(x_b, y_b)} \quad \gamma = 1 - \alpha - \beta$$

- Given a point **p** with Cartesian coordinates (x, y) , we can compute its barycentric coordinates (α, β, γ) as above.

Barycentric Coordinates

- In general, the barycentric area coordinates for point **p** are the solution of the linear system of equations:

$$\begin{pmatrix} x_a & x_b & x_c \\ y_a & y_b & y_c \\ 1 & 1 & 1 \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$



Barycentric Coordinates

- Can be easily converted to trilinear coordinates

$P_t(t_1, t_2, t_3)$ in trilinear coordinates has barycentric coordinates of $(t_1 \mathbf{a}, t_2 \mathbf{b}, t_3 \mathbf{c})$

where $\mathbf{a}, \mathbf{b}, \mathbf{c}$, are the side lengths of the triangle.

$P_b(\alpha, \beta, \gamma)$ in barycentric coordinates has trilinear coordinates $(\alpha/\mathbf{a}, \beta/\mathbf{b}, \gamma/\mathbf{c})$

Triangle Rasterization

- Many different ways to generate fragments for a triangle
- Checking (α, β, γ) is one method, e.g.

$$(0 < \alpha < 1 \ \&\& \ 0 < \beta < 1 \ \&\& \ 0 < \gamma < 1)$$

- In practice, the graphics hardware uses optimized methods:
 - fixed point precision (not floating-point)
 - incremental (use results from previous pixel)

Triangle Rasterization


- We can use barycentric coordinates to rasterize and color triangles

```
for all x do
  for all y do
    compute (alpha, beta, gamma) for (x,y)
    if (0 < alpha < 1 and
        0 < beta < 1 and
        0 < gamma < 1 ) then
      c = alpha c0 + beta c1 + gamma c2
      drawpixel(x,y) with color c
```

- The color c varies smoothly within the triangle

Visibility: One triangle

- With one triangle, things are simple
- Pixels never overlap!




Graphics Lecture 9: Slide 31

Hidden Surface Removal

- Idea: keep track of visible surfaces
- Typically, we see only the front-most surface
- Exception: transparency

Visibility: Two triangles


- Things get more complicated with multiple triangles
- Fragments might overlap in screen space!



Graphics Lecture 9: Slide 33

Visibility: Pixels vs Fragments


- Each pixel has a unique framebuffer (image) location
- But multiple fragments may end up at same address




Graphics Lecture 9: Slide 34

Visibility: Which triangle should be drawn first?

- Two possible cases:




green triangle on top




orange triangle on top

Visibility: Which triangle should be drawn first?

- Many other cases possible!




intersection #1



intersection #2


Visibility: Painter's Algorithm

- Sort triangles (using z values in eye space)
- Draw triangles from back to front



Visibility: Painter's Algorithm - Problems

- Correctness issues:
 - Intersections
 - Cycles
 - Solve by splitting triangles, but ugly and expensive
- Efficiency (sorting)




Graphics Lecture 9: Slide 38

The Depth Buffer (Z-Buffer)

- Perform hidden surface removal per-fragment
- Idea:
 - Each fragment gets a z value in screen space
 - Keep only the fragment with the smallest z value

The Depth Buffer (Z-Buffer)


- Example:
 - fragment from green triangle has z value of 0.7



Graphics Lecture 9: Slide 40

The Depth Buffer (Z-Buffer)


- Example:
 - fragment from red triangle has z value of 0.3



Graphics Lecture 9: Slide 41

The Depth Buffer (Z-Buffer)

- Since $0.3 < 0.7$, the red fragment wins



Graphics Lecture 9: Slide 42

The Depth Buffer (Z-Buffer)

- Many fragments might map to the same pixel location
- How to track their z-values?
- Solution: z-buffer (2D buffer, same size as image)

1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.1	1.0	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.1	0.1	1.0	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.2	0.2	0.3	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.3	0.3	0.4	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.3	0.4	0.4	0.5	1.0	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.4	0.4	0.5	0.5	0.5	1.0	1.0	1.0
1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.4	0.5	1.0	1.0	1.0

Graphics Lecture 9: Slide 43

The Z-Buffer Algorithm

Let CB be color (frame) buffer, ZB be z-buffer

Initialize z-buffer contents to 1.0 (far away)

For each triangle T

 Rasterize T to generate fragments

 For each fragment F with screen position (x,y,z) and color value C

 If (z < ZB[x,y]) then

 Update color: CB[x,y] = C

 Update depth: ZB[x,y] = z


Z-buffer Algorithm Properties

- What makes this method nice?
 - simple (faciliates hardware implementation)
 - handles intersections
 - handles cycles
 - draw opaque polygons in any order


Alias Effects

- One major problem with rasterization is called alias effects, e.g straight lines or triangle boundaries look jagged
- These are caused by undersampling, and can cause unreal visual artefacts.
- It also occurs in texture mapping

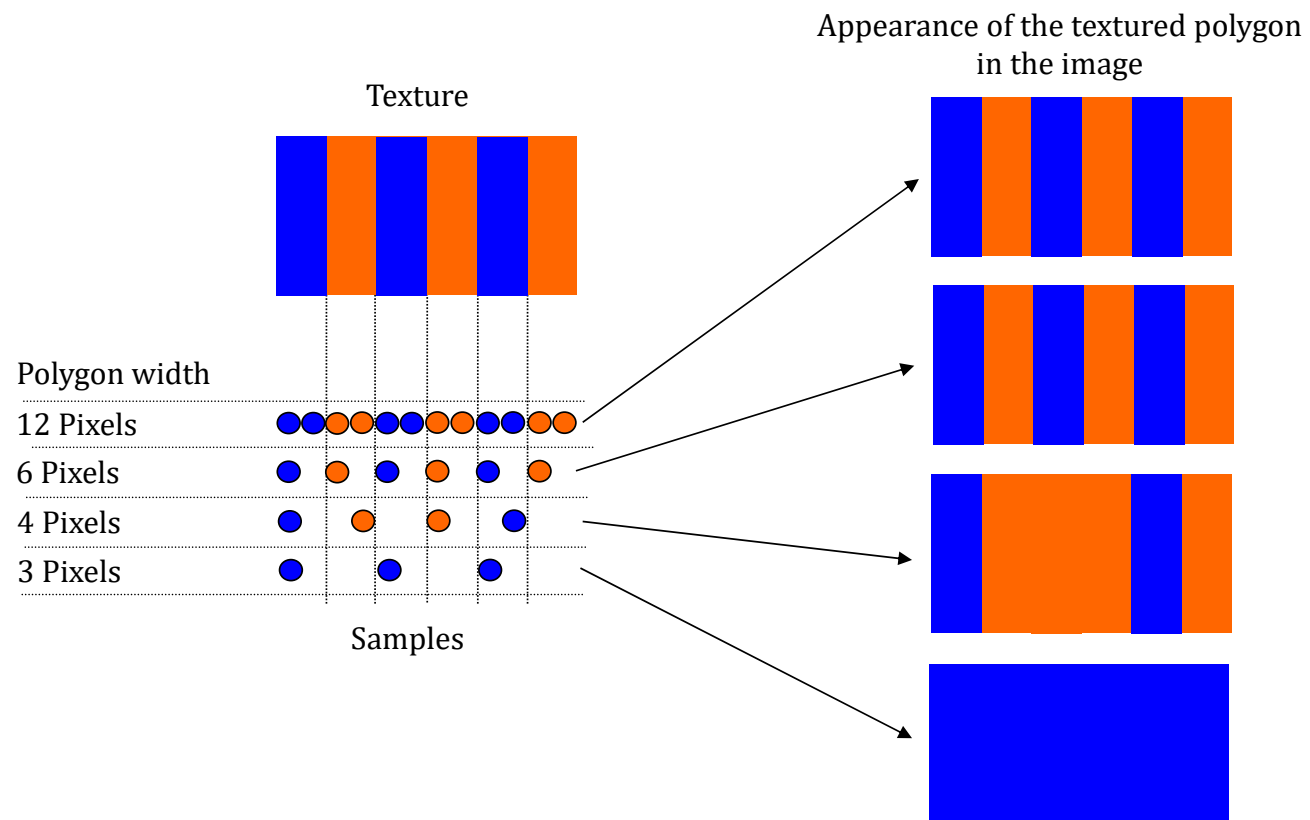
Alias Effects at straight boundaries in raster images.



Desired Boundaries



Pixels Set




Anti-Aliasing

- The solution to aliasing problems is to apply a degree of blurring to the boundary such that the effect is reduced.
- The most successful technique is called **Supersampling**



Supersampling

- The basic idea is to compute the picture at a higher resolution to that of the display area.
- Supersamples are averaged to find the pixel value.
- This has the effect of blurring boundaries, but leaving coherent areas of colour unchanged



Limitations of Supersampling


- Supersampling works well for scenes made up of filled polygons.
- However, it does require a lot of extra computation.
- It does not work for line drawings.




Convolution filtering

- The more common (and much faster) way of dealing with alias effects is to use a 'filter' to blur the image.
- This essentially takes an average over a small region around each pixel

For example consider the image of a line



Treat each pixel of the image





We replace the pixel by a local average,
one possibility would be $3 \cdot I/9$


Weighted averages

- Taking a straight local average has undesirable effects.
- Thus we normally use a weighted average.

$$1/36 *$$

1	4	1
4	16	4
1	4	1





Pros and Cons of Convolution filtering

- Advantages:
 - It is very fast and can be done in hardware
 - Generally applicable
- Disadvantages:
 - It does degrade the image while enhancing its visual appearance.

Anti-Aliasing textures

- Similar
- When we identify a point in the texture map we return an average of texture map around the point.
- Scaling needs to be applied so that the less the samples taken the bigger the local area where averaging is done.

